

# Basic Principles

---

## Big Data Management

Anis Ur Rahman

Faculty of Computer Science & Information Technology  
University of Malaya

October 15, 2019

# Lecture Outline

## Different aspects of data distribution

### 1 Scaling

- Vertical vs. horizontal

### 2 Distribution models

- Sharding
- Replication: master-slave vs. peer-to-peer architectures

### 3 CAP properties

- Consistency, availability and partition tolerance
- ACID vs. BASE guarantees

### 4 Consistency

- Read and write quora

# Outline

- 1 Scalability
- 2 Distribution Models
- 3 CAP Theorem
- 4 Consistency

# Scalability

## What is scalability?

- the capability of a system **to handle growing** amounts of **data** and/or **queries** without losing performance, or
  - its potential to be **enlarged** in order to accommodate such a growth

## Two general approaches


- **Vertical** scaling
- **Horizontal** scaling

# Vertical Scalability

## Vertical scaling (scaling up/down)

- Adding resources to a **single node** in a system
  - e.g. increasing the number of CPUs, extending system memory, using larger disk arrays, ...
  - i.e. **larger and more powerful** machines are involved
- Traditional choice
  - In favor of **strong** consistency
  - **Easy** to implement and deploy
  - **No** issues caused by data distribution
  - ...

Works well in many cases but ...



# Vertical Scalability: Drawbacks

## Performance limits

- Even the most powerful machine has a **limit**
- Everything works well ... **until** start approaching the limit

## Higher costs

- The cost of expansion **increases exponentially**
- In particular, it is higher than the sum of costs of equivalent commodity hardware

## Proactive provisioning

- New projects/applications might **evolve rapidly**
- **Upfront budget** is needed when deploying new machines
- So **flexibility** is seriously suppressed

# Vertical Scalability: Drawbacks

## Vendor lock-in

- There are only a few manufacturers of large machines
- Customer is made **dependent** on a **single vendor**
  - Their products, services, but also implementation details, proprietary formats, interfaces, support, ...
- e.e. it is **difficult** or **impossible** to **switch** to another vendor

## Deployment downtime

- **Inevitable** downtime is often required when scaling up


# Horizontal Scalability

## Horizontal scaling (scaling out/in)

- **adding** more nodes to a system
  - i.e. system is **distributed** across multiple nodes in a cluster
- **Choice** of many **NoSQL** systems

## Advantages

- Commodity hardware, **cost effective**
- **Flexible** deployment and maintenance
- Often **surpasses** the vertical scaling
- Often **no** single point of failure
- ...





# Horizontal Scalability: Consequences

## Significantly increases complexity

- Complexity of management, programming model, ...

## Introduces new issues and problems

- Data **distribution**
- **Synchronization** of nodes
- Data **consistency**
- **Recovery** from failures
- ...

And there are also plenty of false assumptions ...

# Horizontal Scalability: Fallacies

## False assumptions

- Network is **reliable**
- Latency is **zero**
- Bandwidth is **infinite**
- Network is **secure**
- Topology does **not change**
- There is **one** administrator
- Network is **homogeneous**
- Transport cost is **zero**

Source: <https://www.red-gate.com/simple-talk/blogs/the-eight-fallacies-of-distributed-computing/>

## Horizontal Scalability: Conclusion

A **standalone node** still might be a better option in certain cases

- e.g. for graph databases
  - Simply because it is difficult to **split** and **distribute** graphs
- In other words
  - It can make sense to run even a **NoSQL** database system on a **single node**
  - **No distribution** at all is the most **preferred/simple** scenario

But in general, horizontal scaling really opens **new possibilities**

# Horizontal Scalability: Architecture

## What is a cluster?

- A **collection** of mutually interconnected commodity nodes
- Based on the **shared-nothing** architecture
  - Nodes **do not share** their CPUs, memory, hard drives, ...
  - Each node **runs its own** operating system instance
  - Nodes **send messages** to interact with each other
- Nodes of a cluster can be **heterogeneous**
- Data, queries, calculations, requests, workload, ... all **distributed** among the nodes within a cluster

# Outline

- 1 Scalability
- 2 Distribution Models**
- 3 CAP Theorem
- 4 Consistency

# Distribution Models

## Generic techniques of data distribution

### 1 Sharding

- Idea. **different data** on different nodes
- Motivation. increasing **volume** of data, increasing **performance**

### 2 Replication

- Idea. the **same data** on different nodes
- Motivation. increasing **performance**, increasing **fault tolerance**

Both the techniques are **mutually orthogonal**

- i.e. we can **use either** of them, or **combine** them both

## Distribution model

- specific way how sharding and replication is **implemented**

NoSQL systems often offer **automatic** sharding and replication

# Sharding


## Sharding (horizontal partitioning)

- **Placement** of different data on different nodes
  - What different data means? Usually **aggregates**
    - e.g. key-value pairs, documents, ...
- **Related pieces** of data that are **accessed together** should also be **kept together**
  - Specifically, operations involving data on **multiple shards** should be **avoided** (if possible)

The questions are...

- how to **design** aggregate structures?
- how to actually **distribute** these aggregates?

# Sharding



Source: Sadalage, Pramod J. - Fowler, Martin: NoSQL Distilled. Pearson Education, Inc., 2013.



# Sharding

## Objectives

- Achieve **uniform** data distribution
- Achieve **balanced** workload (read and write requests)
- Respect physical **locations**
  - e.g. different data centers for users around the world
- ...

Unfortunately, these objectives...

- may **mutually contradict** each other
- may **change** in time

So, how to actually determine shards for aggregates?

# Sharding

Original Table

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN
3	SELDA	BAGCAN	PURPLE
4	JIM	PEPPER	AUBERGINE

Vertical Partitions

VP1

CUSTOMER ID	FIRST NAME	LAST NAME
1	TAEKO	OHNUKI
2	O.V.	WRIGHT
3	SELDA	BAGCAN
4	JIM	PEPPER

VP2

CUSTOMER ID	FAVORITE COLOR
1	BLUE
2	GREEN
3	PURPLE
4	AUBERGINE

Horizontal Partitions

HP1

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN

HP2

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
3	SELDA	BAGCAN	PURPLE
4	JIM	PEPPER	AUBERGINE

Source:

<https://www.digitalocean.com/community/tutorials/understanding-database-sharding>

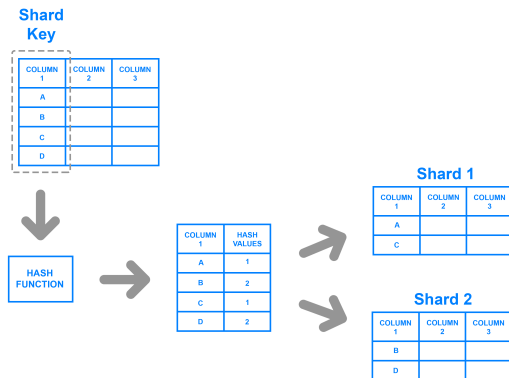
# Sharding

## Sharding strategies

- Based on **mapping structures**
  - Data is placed on shards in a **random** fashion
  - e.g. round-robin, ...
- Knowledge of the **mapping** of individual aggregates to particular shards must then be **maintained**
  - Thus usually maintained using a **centralized index structures** with all the disadvantages
- Based on general rules
  - Each shard is responsible for **storing certain data**
  - Hash partitioning, range partitioning, ...

# Sharding

## Key-based sharding

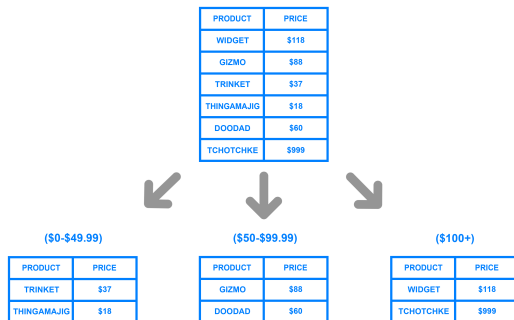


Source:

<https://www.digitalocean.com/community/tutorials/understanding-database-sharding>

# Sharding

## Range-based sharding

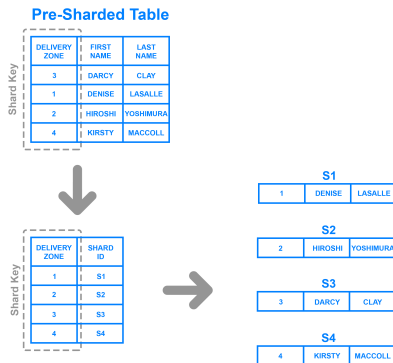


Source:

<https://www.digitalocean.com/community/tutorials/understanding-database-sharding>

# Sharding

## Directory-based sharding



Source:

<https://www.digitalocean.com/community/tutorials/understanding-database-sharding>

# Sharding

## Should I Shard?

- Amount of **application data** grows to exceed the storage capacity of a single database node.
- Volume of **writes** or **reads** to the database surpasses what a single node can handle,
  - resulting in slowed response times or timeouts.

# Sharding

## Why is sharding difficult?

- Not only we need to be able to determine particular shards during **write requests**
  - i.e. when a new aggregate is about to be **inserted**
  - So that we can actually make a **decision** where it should be physically stored
- but also during **read requests**
  - i.e. when existing aggregate/s are about to be retrieved
  - So that we can actually **find** and **return** them efficiently (or detect they are missing)
  - And all that only based on the **search criteria** provided (e.g. key, id, ...) unless all the nodes should be accessed



# Sharding

## Why is sharding even more difficult?

- **Structure** of the cluster may be **changing**
  - Nodes can be **added** or **removed**
- Nodes may have **incomplete/obsolete** cluster knowledge
  - Nodes involved, their responsibilities, sharding rules, ...
- **Individual nodes** may be **failing**
- **Network** may be **partitioned**
  - Messages may not be delivered even though sent

# Replication

## Replication


- Placement of **multiple copies** of the same data (replicas) on different nodes
- **Replication factor** = number of such copies

## Two approaches

- 1 Master-slave architecture
- 2 Peer-to-peer architecture

# Replication

## Master-Slave Architecture



Source: Sadalage, Pramod J. - Fowler, Martin: NoSQL Distilled. Pearson Education, Inc., 2013.

# Replication

## Master-Slave Architecture

- One node is **primary** (master), all the other **secondary** (slave)
- Master node bears all the **management** responsibility
- All the nodes contain **identical** data
  - **Read requests** can be handled by **both** the master or slaves
- Suitable for **read-intensive** applications
  - **More read requests** to deal with → **more slaves** to deploy
- When the **master fails**, read operations can still be **handled**

# Replication


## Master-Slave Architecture

Write requests can **only** be handled by the master

- Newly written **replicas** are **propagated** to all the slaves
- **Consistency issue**
  - Luckily enough, **at most one** write request is handled at a time
  - But the propagation still **takes some time** during which **obsolete reads** might happen
  - Hence certain **synchronization** is required **to avoid conflicts**
- In case of **master failure**, a new one needs to be **appointed**
  - **Manually** (user-defined) or **automatically** (cluster-elected)
  - Since the nodes are **identical**, appointment can be **fast**
- Master might therefore represent a **bottleneck** (because of the performance or failures)

# Replication

## Peer-to-Peer Architecture



Source: Sadalage, Pramod J. - Fowler, Martin: NoSQL Distilled. Pearson Education, Inc., 2013.

# Replication

## Peer-to-Peer Architecture

- All the nodes have **equal** roles and responsibilities
- All the nodes contain **identical data** once again
  - **Both** read and write requests can be handled by **any node**
- No **bottleneck**, no single point of **failure**
- Both the operations **scale** well
  - **More requests** to deal with → **more nodes** to deploy
- **Consistency issues**
  - Unfortunately, **multiple** write requests can be initiated **independently** and being executed at the **same time**
  - Hence **synchronization** is required to avoid conflicts

# Sharding and Replication

Observations with respect to the replication

- Does the **replication factor** really need to correspond to the number of nodes?
  - No, replication factor of 3 will often be the right choice
  - **Consequences**
    - Nodes will **no** longer contain **identical data**
    - **Replica placement** strategy will be needed
- Do all the replicas really need to be **successfully written** when write requests are handled?
  - **No**, but consistency issues have to be tackled carefully

Sharding and replication can be combined... but how?



# Sharding and Replication

## Sharding and Master-Slave Replication

master for two shards



slave for two shards



master for one shard



master for one shard  
and slave for a shard



slave for two shards




slave for one shard



Source: Sadalage, Pramod J. - Fowler, Martin: NoSQL Distilled. Pearson Education, Inc., 2013.

# Sharding and Replication

## Sharding and Peer-to-Peer Replication



Source: Sadalage, Pramod J. - Fowler, Martin: NoSQL Distilled. Pearson Education, Inc., 2013.

# Sharding and Replication

## Combinations of sharding and replication

### 1 Sharding + master-slave replication

- **Multiple masters**, each for different data
- **Roles** of the nodes can **overlap**
  - Each node can be **master** for some data and/or **slave** for other

### 2 Sharding + peer-to-peer replication

- Basically **placement of anything anywhere** (although certain rules can still be applied)

# Sharding and Replication

Questions to figure out for any distribution model

- Can all the nodes **serve both** read and write requests?
- Which **replica placement strategy** is used?
- How the **mapping of replicas** is maintained?
- What **level of consistency** and **availability** is provided?
- What extent of **infrastructure knowledge** do the nodes have?
- ...

# Outline

- 1 Scalability
- 2 Distribution Models
- 3 CAP Theorem**
- 4 Consistency

# CAP Theorem

## Assumptions

- **Distributed** system with sharding and replication
- Read and write operations on a **single aggregate only**

## CAP properties

- Properties of a distributed system:
- Consistency, Availability, and Partition tolerance

## CAP theorem

- It is **not possible** to have a distributed system that would **guarantee all** CAP properties at the same time.
- Only 2 of these 3 properties can be enforced.

But, what these properties actually mean?

# CAP Properties

## Consistency

- Read and write **operations** must be executed **atomically**
  - There must exist a total **order on all operations**
  - Each operation looks as if it was **completed at a single instant**
  - i.e. as if all operations were executed **sequentially** one by one on a single standalone node
- **Practical consequence.** after a write operation, all readers **see the same data**
  - Since any node can be used for handling of read requests, atomicity of write operations means that **changes must be propagated** to all the replicas
  - As we will see later on, other ways for such a strong consistency exist as well

# CAP Properties

## Availability

- If a node is working, it **must respond** to user requests
  - Every read or write request successfully received by a non-failing node in the system must result in a response,
  - i.e. their execution must not be rejected

## Partition tolerance

- System **continues to operate** even when two or more sets of **nodes get isolated**
- The network is **allowed to lose** arbitrarily many messages sent from one node to another
  - i.e. a connection failure **must not shut** the whole system down



# CAP Theorem Consequences

If at most two properties can be **guaranteed** ...

## 1 CA = consistency + availability

- Traditional **ACID** properties are **easy to achieve**
- **Examples.:** RDBMS, Google BigTable
- Any single-node system, but even clusters (at least in theory)
  - However, should the **network partition** happen, all the nodes must be **forced to stop** accepting user requests


## 2 CP = consistency + partition tolerance

- **Other examples.** distributed locking

## 3 AP = availability + partition tolerance

- New concept of **BASE** properties
- **Examples.** Apache Cassandra, Apache CouchDB
- **Other examples.** web caching, DNS

# CAP Theorem Consequences



# CAP Theorem Consequences

## Partition tolerance is necessary in clusters

- Why? Because it is **difficult to detect** network failures
- Does it mean that only purely CP and AP systems are possible?
- No...

## The real meaning of the CAP theorem:

- The real-world does not need to be just black and white
- **Partition tolerance** is a **must**, but we can **trade off** consistency versus availability
  - Just a little bit **relaxed consistency** can bring a lot of availability
  - Such trade-offs are not only possible, but often **works very well** in practice

## CAP Theorem Example

You are asked to design a distributed cluster of 4 data nodes. Replication factor is 2 i.e. any data written in cluster must be written on 2 nodes; so when one goes down – second can serve the data. Now try to apply CAP theorem on this requirement.

In distributed system, two things may happen anytime i.e. node failure (hard disk crash) or network failure (connection between two nodes go down).

- 1 CP [Consistency/Partition Tolerance] Systems
- 2 AP [Availability/Partition Tolerance] Systems
- 3 CA [Consistency/Availability] Systems

# ACID Properties

## Traditional ACID properties

### 1 Atomicity

- **Partial** execution of transactions is **not allowed** (all or nothing)

### 2 Consistency

- Transactions bring the database from **one consistent (valid) state to another**

### 3 Isolation

- Transactions executed in parallel **do not see uncommitted effects** of each other

### 4 Durability

- Effects of **committed transactions** must remain **durable**

# BASE Properties

## New concept of BASE properties

### 1 Basically Available

- The system **works** basically **all the time**
- **Partial failures** can occur, but there are **no total system failures**

### 2 Soft State

- The system is in **flux (unstable)**, **non-deterministic state**
- **Changes** occur all the time

### 3 Eventual Consistency

- **Sooner or later** the system will be in **some consistent state**

BASE is just a vague term, no formal definition

- Proposed to illustrate design philosophies at the opposite ends of the consistency-availability spectrum

# ACID and BASE

## ACID

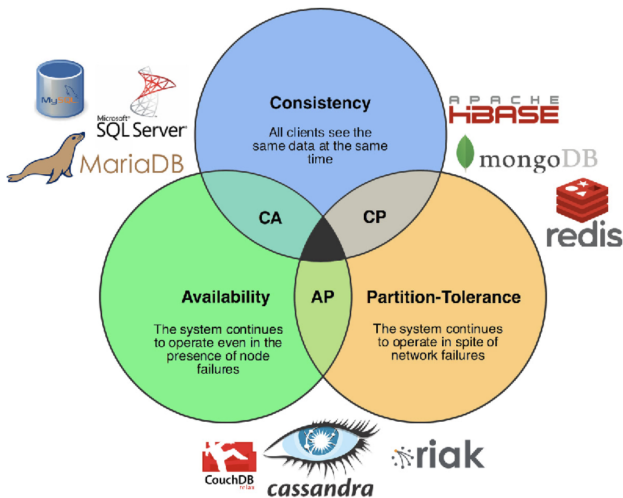
- Choose **consistency over availability**
- **Pessimistic** approach
- Implemented by **traditional** relational databases

## BASE

- Choose **availability over consistency**
- **Optimistic** approach
- Common in **NoSQL databases**
- Allows levels of **scalability** that cannot be acquired with ACID

**Current trend in NoSQL.** strong consistency → eventual consistency

# CAP Theorem Conclusion





# Outline

- 1 Scalability
- 2 Distribution Models
- 3 CAP Theorem
- 4 Consistency**

# Consistency

## Consistency in general...

- Consistency is the **lack of contradiction** in the database
- However, it has many facets ... e.g.
  - only assume **atomic operations** always **manipulating** just a **single aggregate**,
  - but **set operations** could also be considered etc.

**Strong consistency** is **achievable** even in clusters, but **eventual consistency** might often be **sufficient**

- 1 One minute obsolete article on a news portal does not matter
- 2 Even when an already unavailable hotel room is booked once again, the situation can still be figured out in the real world
- 3 ...

# Consistency

## Write consistency (update consistency)

- Problem. **write-write conflict**
  - Two or more **write requests** on the same aggregate are **initiated concurrently**
- Context. **peer-to-peer** architecture only
- Issue. **lost update**
- Solution.
  - 1 **Pessimistic strategies**
    - Preventing conflicts from occurring
    - Write locks, ...
  - 2 **Optimistic strategies**
    - Conflicts may occur, but are **detected** and **resolved** later on
    - Version stamps, vector clocks, ...

# Consistency

## Read consistency (replication consistency)

- Problem. **read-write conflict**
  - Write and read requests on the same aggregate are **initiated concurrently**
- Context. **both** master-slave and peer-to-peer architectures
- Issue. **inconsistent read**
- When not treated, **inconsistency window** will exist
  - **Propagation** of changes to all the replicas **takes some time**
  - Until this process is finished, inconsistent reads may happen
  - Even the initiator of the write request may read wrong data!
    - Session consistency/read-your-writes/sticky session

# Strong Consistency

How many nodes need to be involved to get strong consistency?

- **Write quorum.**  $W > N/2$ 
  - **Idea.** only one write request can get the **majority**
  - $W$  = number of nodes successfully participating in the write
  - $N$  = number of nodes involved in replication (replication factor)
- **Read quorum.**  $R > N - W$ 
  - **Idea.** concurrent write requests cannot happen
  - $R$  = number of nodes participating in the read
  - Should the retrieved replicas be mutually different, the newest version is resolved and then returned

When a quorum is not attained *rightarrow* the request cannot be handled

# Strong Consistency

Examples for replication factor  $N = 3$

- Write quorum  $W = 3$  and read quorum  $R = 1$ 
  - **All** the replicas are always **updated**
  - Can read **any** one of them
- Write quorum  $W = 2$  and read quorum  $R = 2$ 
  - **Typical** configuration, **reasonable** trade-off

Consequence

- Quora can be **configured** to **balance** read and write workload
  - The **higher** the write quorum is required **lower** the read quorum can then be required

## Lecture Conclusion

There is a **wide range of options** influencing...

- **Scalability** – how well the entire system scales?
- **Availability** – when nodes may refuse to handle user requests?
- **Consistency** – what level of consistency is required?
- **Latency** – how long does it take to handle user requests?
- **Durability** – is the committed data written reliably?
- **Resilience** – can the data be recovered in case of failures?

it's good to know these properties and choose the right trade-off