

Document Databases: MongoDB

Big Data Management

Faculty of Computer Science & Information Technology
University of Malaya

December 3, 2019

Lecture Outline

Document databases

- Introduction

MongoDB

- Data model
- CRUD operations
 - Insert
 - Update
 - Remove
 - Find: projection, selection, modifiers

Outline

- 1 Background
- 2 MongoDB Document Database
- 3 Insert Operation
- 4 Update Operation
- 5 Remove Operation
- 6 Find Operation

Functionality of MongoDB

- Dynamic schema
 - No DDL
- Document-based database
- Secondary indexes
- Query language via an API
- Atomic writes and fully-consistent reads
 - If system configured that way
- Master-slave replication with automated failover (replica sets)
- Built-in horizontal scaling via automated range-based partitioning of data (sharding)
- No joins nor transactions

Why use MongoDB?

- Simple queries
- Functionality provided applicable to most web applications
- Easy and fast integration of data
 - No ERD diagram
- Not well suited for heavy and complex transactions systems

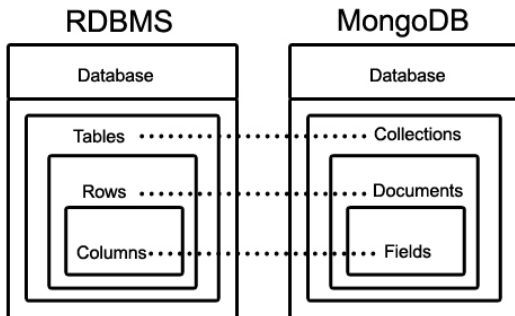
MongoDB: CAP approach

Focus on Consistency and Partition tolerance

- **Consistency**
 - all replicas contain the same version of the data
- **Availability**
 - system remains operational on failing nodes
- **Partition tolerance**
 - multiple entry points
 - system remains operational on system split

MongoDB: Hierarchical Objects

- A MongoDB instance may have zero or more **databases**
- A database may have zero or more **collections**
- A collection may have zero or more **documents**
- A document may have one or more **fields**
- MongoDB **Indexes** function much like their RDBMS counterparts.



RDB Concepts to NO SQL

- Database \Rightarrow Database
- Table, View \Rightarrow Collection
- Row \Rightarrow Document (BSON)
- Column \Rightarrow Field
- Index \Rightarrow Index
- Join \Rightarrow Embedded Document
- Foreign Key \Rightarrow Reference
- Partition \Rightarrow Shard

Collection is not strict about what it stores

- Schema-less

Hierarchy is evident in the design

- Embedded Document

MongoDB Processes and configuration

- **Mongod** - Database instance
- **Mongos** - Sharding processes
 - Analogous to a **database router**.
 - **Processes** all requests
 - **Decides** how many and which mongods should receive the query
 - **Collates** the results, and sends it back to the client.
- **Mongo** – an interactive shell (a client)
 - Fully functional JavaScript environment for use with a MongoDB
- You can have **one mongos** for the whole system no matter how many mongods you have
- OR you can have **one local mongos** for every client if you wanted to minimize network latency.

Choices made for Design of MongoDB

- **Scale horizontally** over commodity hardware
 - Lots of relatively inexpensive servers
- **Keep the functionality** that works well in RDBMSs
 - Ad hoc queries
 - Fully featured indexes
 - Secondary indexes
- What doesn't distribute well in RDB?
 - Long running multi-row transactions
 - Joins
 - Both artifacts of the relational data model (row x column)

BSON format

- Binary-encoded serialization of JSON-like documents
- Zero or more key/value pairs are stored as a single entity
- Each entry consists of a field name, a data type, and a value
- Large elements in a BSON document are prefixed with a length field to facilitate scanning

Schema Free

- MongoDB does **not** need any **pre-defined data schema**
- Every document in a collection could have **different data**
 - Addresses **NULL** data fields

JSON format

Data is in name/value pairs

- A name/value pair consists of a field name followed by a colon, followed by a value:
 - Example: 'name': 'R2-D2'
- Data is separated by commas
 - Example: 'name': 'R2-D2', race : 'Droid'
- Curly braces hold objects
 - Example: 'name': 'R2-D2', race : 'Droid', affiliation: 'rebels'
- An array is stored in brackets []
 - Example ['name': 'R2-D2', race : 'Droid', affiliation: 'rebels', 'name': 'Yoda', affiliation: 'rebels']

Document Stores

Data model

- **Documents**
 - Self-describing
 - **Hierarchical** tree structures (JSON, XML, ...)
 - Scalar values, maps, lists, sets, nested documents, ...
 - Identified by a **unique identifier** (key, ...)
- Documents are organized into **collections**

Query patterns

- Create, update or remove a document
- Retrieve documents according to complex query conditions

Observation

- **Extended key-value stores** where the value part is **examinable**

Outline

- 1 Background
- 2 MongoDB Document Database**
- 3 Insert Operation
- 4 Update Operation
- 5 Remove Operation
- 6 Find Operation

MongoDB

JSON document database

- <https://www.mongodb.com/>
- **Features**
 - Open source, high availability, eventual consistency, automatic sharding, master-slave replication, automatic failover, secondary indices, ...
- Developed by MongoDB
- Implemented in C++, C, and JavaScript
- Operating systems: Windows, Linux, Mac OS X, ...
- Initial release in **2009**

Query Example

Collection of movies

```
1 {  
2   _id: ObjectId("1"),  
3   title: "The Double",  
4   year: 2006  
5 }
```

```
1 {  
2   _id: ObjectId("2"),  
3   title: "Zombieland",  
4   year: 2000  
5 }
```

```
1 {  
2   _id: ObjectId("3"),  
3   title: "The Social Network",  
4   year: 2007  
5 }
```

Query statement

Titles of movies filmed in 2005 and later, sorted by these titles in descending order

```
1 db.movies.find(  
2   { year: { $gt: 2005 } },  
3   { _id: false, title: true }  
4 ).sort([ title: -1 ])
```

Query result

```
1 { title: "The Double" }
```

```
1 { title: "The Social Network" }
```

Data Model

Database system structure

Instance → databases → collections → documents

1 Database

2 Collection

- Collection of documents, usually of a **similar structure**

3 Document

- MongoDB document = **one JSON object**
 - i.e. even a complex JSON object with other recursively nested objects, arrays or values
- Each document has a **unique identifier** (primary key)
 - Technically realized using a **top-level _id field**

Data Model

MongoDB document

- Internally stored in **BSON format** (Binary JSON)
 - Maximal allowed size **16 MB**
 - **GridFS** can be used to split larger files into smaller chunks

Restrictions on field names

- Top-level `_id` is **reserved** for a primary key
- Field names **cannot** start with `$`
 - Reserved for **query operators**
- Field names **cannot** contain `.`
 - Used when **accessing nested fields**

Primary Keys

Features of identifiers

- **Unique** within a collection
- **Immutable** (cannot be changed once assigned)
- Can be of **any type** other than a JSON array

Key management

- 1 **Natural** identifier
- 2 **Auto-incrementing** number – not recommended
- 3 **UUID** (Universally Unique Identifier)
- 4 **ObjectId** – special 12-byte BSON type (the default option)
 - Small, likely unique, fast to generate, ordered, based on a timestamp, machine id, process id, and a process-local counter

Design Questions

Data modeling (in terms of collections and documents)

- **No explicit schema** is provided, nor expected or enforced
 - However...
 - documents within a collection are **similar in practice**
 - **implicit schema** is required nevertheless
- **Challenge**
 - Balancing application requirements, performance aspects, data structure, mutual relationships, query patterns, ...

Two main concepts

- 1 Embedded documents
- 2 References

Denormalized Data Models

Embedded documents

- **Related** data in a **single document**
 - with **embedded JSON objects**, so called **subdocuments**
- **Pros:** data **manipulation** (fewer queries need to be issued)
- **Cons:** possible data **redundancies**
- Suitable for **one-to-one** or **one-to-many** relationships

```
1 {  
2   _id: ObjectId("2"), title: "Zombieland", year: 2000,  
3   actors: [  
4     { firstname: "Emma", lastname: "Stone" },  
5     { firstname: "Woody", lastname: "Harrelson" },  
6     { firstname: "Jesse", lastname: "Eisenberg" }  
7   ]  
8 }
```

Normalized Data Models

References

- **Related data in separate documents**
 - These are interconnected via **directed links** (references)
 - Technically expressed using ordinary values with **identifiers of target documents** (i.e. no special construct is provided)
- **Features:** higher flexibility, follow up queries might be needed
- Suitable for **many-to-many** relationships

```
1 {  
2   _id: ObjectId("2"),  
3   title: "Zombieland",  
4   year: 2000,  
5   actors: [  
6     ObjectId("6"),  
7     ObjectId("4"),  
8     ObjectId("5") ]  
9 }
```

```
1 {  
2   _id: ObjectId("6"),  
3   firstname: "Emma",  
4   lastname: "Stone"  
5 }
```

Sample Data

Collection of movies

```

1 {
2   _id: ObjectId("1"),
3   title: "The Double", year: 2006,
4   actors: [ ObjectId("7"), ObjectId("5") ]
5 }

```

```

1 {
2   _id: ObjectId("2"),
3   title: "Zombieland", year: 2000,
4   actors: [ ObjectId("6"), ObjectId("4"), ←
5             ObjectId("5") ]

```

```

1 {
2   _id: ObjectId("3"),
3   title: "The Social Network", year: 2007,
4   actors: [ ObjectId("5"), ObjectId("4") ]
5 }

```

Collection of actors

```

1 { _id: ObjectId("4"),
2   firstname: "Woody",
3   lastname: "Harrelson" }

```

```

1 { _id: ObjectId("5"),
2   firstname: "Jesse",
3   lastname: "Eisenberg" }

```

```

1 { _id: ObjectId("6"),
2   firstname: "Emma",
3   lastname: "Stone" }

```

```

1 { _id: ObjectId("7"),
2   firstname: "Mia",
3   lastname: "Wasikowska" }

```


Application Interfaces

mongo shell

- Interactive interface to MongoDB
- `./bin/mongo --username user --password pass --host host --port 28015`

Drivers

- Java, C, C++, C, Perl, PHP, Python, Ruby, Scala, ...

Query Language

MongoDB query language is based on **JavaScript**

- Single command/entire script
- Read queries return a cursor
 - Allows us to **iterate** over all the selected documents
- Each command is always **evaluated** over a **single collection**

Query patterns

- Basic CRUD operations
 - Accessing documents **via identifiers** or **conditions on fields**
- **Aggregations:** MapReduce, pipelines, grouping

CRUD Operations

Overview


- 1 `db.collection.insert()`
 - **Inserts** a new document into a collection
- 2 `db.collection.update()`
 - **Modifies** an existing document/documents or inserts a new one
- 3 `db.collection.remove()`
 - **Deletes** an existing document/documents
- 4 `db.collection.find()`
 - **Finds** documents based on filtering conditions
 - **Projection** and/or **sorting** may be applied too

Outline

- 1 Background
- 2 MongoDB Document Database
- 3 Insert Operation**
- 4 Update Operation
- 5 Remove Operation
- 6 Find Operation

Insert Operation

Inserts a new document/documents into a given collection



- Parameters

- 1 **Document:** **one or more** documents to be inserted
 - Provided document identifiers (`_id` fields) must be **unique**
 - When **missing**, they are **generated automatically** (ObjectId)

2 Options

- Collections are **created automatically** when not yet exist

Insert Operation: Examples

Insert a new actor document

```
1 db.actors.insert(  
2   {  
3     firstname: "Jennifer",  
4     lastname: "Lawrence"  
5   }  
6 )
```

```
1 {  
2   _id: ObjectId("8"),  
3   firstname: "Jennifer",  
4   lastname: "Lawrence"  
5 }
```

Insert two new movies


```
1 db.movies.insert(  
2   [  
3     {  
4       _id: ObjectId("9"), title: "The Hunger Games", year: 2003,  
5       actors: [ ObjectId("4"), ObjectId("8") ]  
6     },  
7     { title: "Passengers", year: 2016, actors: [ ObjectId("8") ] },  
8   ]  
9 )
```

Outline

- 1 Background
- 2 MongoDB Document Database
- 3 Insert Operation
- 4 Update Operation**
- 5 Remove Operation
- 6 Find Operation

Update Operation

Modifies/replaces an existing document/documents



Parameters

- 1 **Query:** description of documents to be updated
 - The **same** behavior as in find operations
 - 2 **Update:** modification actions to be applied
 - 3 **Options**
- **At most one document** is updated by default
 - Unless **{ multi: true }** option is specified

Update Operation: Examples

Replace the whole document of at most one specified actor

```
1 db.actors.update(  
2   { _id: ObjectId("8") },  
3   {   firstname: "Ana",  
4       lastname: "Lawrence" }  
5 )
```


```
1 {  
2   _id: ObjectId("8"),  
3   firstname: "Ana",  
4   lastname: "Lawrence"  
5 }
```

Update all movies filmed in 2015 or later

```
1 db.movies.update(  
2   { year: { $gt: 2015 } },  
3   {  
4     $set: { new: true },  
5     $inc: { rating: 3 }  
6   },  
7   { multi: true }  
8 )
```

Update Operation

Update/replace modes



1 Replace

- when the update parameter contains **no update operators**
- The **whole document is replaced** (`_id` is preserved)


2 Update

- when the update parameter contains **only update operators**
- **Current document is updated** using these operators
 - `$set`, `$unset`, `$inc`, `$mul`, ...
 - Each operator can be used **at most once**


Update Operators

Field operators


- 1 **\$set** – sets the value of a given field/fields



- 2 **\$unset** – removes a given field/fields




- 3 **\$rename** – renames a given field/fields




Update Operators

Field operators


- 1 **\$inc** – increments the value of a given field/fields



- 2 **\$mul** – multiplies the value of a given field/fields




- 3 **\$currentDate** – stores the current date time/timestamp to a given field/fields




Update Operators

Array operators

- 1 **\$push** – adds one item/all items to the end of an array




- 2 **\$addToSet** – adds one item/all items to the end of an array, but duplicate values are ignored




Update Operators

Array operators

- 1 **\$pop** – removes the first/last item of an array



- 2 **\$pull** – removes all array items that match a specified query



Upsert Mode

Upsert behavior of update operation

- When { **upsert: true** } option is **specified**, and,
- **no document was updated** \Rightarrow new document is **inserted**

What this document will contain?

- In case of the **replace mode**...
 - All the **fields** (i.e. value fields) from the **update parameter**
- In case of the **update mode**...
 - All the **value fields** from the **query parameter**, and
 - All the **outcome** of all the **update operators** from the update parameter
- **_id** field is **preserved**, or **newly generated** if necessary

Upsert Mode: Example

Unsuccessful update of a movie resulting to an insertion

```
1 db.movies.update(  
2   { title: "Joy", year: { $gt: 2000 } },  
3   {  
4     $set: {  
5       director: { firstname: "David", lastname: "Russell" },  
6       year: 2001  
7     },  
8     $inc: { rating: 2 }  
9   },  
10  { upsert: true }  
11 )
```


```
1 { _id: ObjectId("11"),  
2   title: "Joy",  
3   director: { firstname: "David", lastname: "Russell" },  
4   year: 2001,  
5   rating: 2  
6 }
```


Outline

- 1 Background
- 2 MongoDB Document Database
- 3 Insert Operation
- 4 Update Operation
- 5 Remove Operation**
- 6 Find Operation

Remove Operation

Removes a document/documents from a given collection



- **Parameters**

- 1 **Query:** description of documents to be removed

- The **same behavior** as in find operations

- 2 **Options**


- All the matching documents are **removed unless** `{ justOne: true }` option is provided

Outline

- 1 Background
- 2 MongoDB Document Database
- 3 Insert Operation
- 4 Update Operation
- 5 Remove Operation
- 6 Find Operation**

Find Operation

Selects documents from a given collection



- Parameters

- 1 **Query:** description of documents to be **selected**
- 2 **Projection:** fields to be **included/excluded** in the result

- Matching documents are returned via an **iterable cursor**
 - This allows us to chain further **sort**, **skip** or **limit** operations

Find Operation: Examples

Select all movies from our collection

```
1 db.movies.find()
```

```
1 db.movies.find( { } )
```

Select a particular movie based on its document identifier

```
1 db.movies.find( { _id: ObjectId("2") } )
```

Select movies filmed in 2000 with a rating greater than 1


```
1 db.movies.find( { year: 2000, rating: { $gt: 1 } } )
```

Select movies filmed between 2005 and 2015

```
1 db.movies.find( { year: { $gte: 2005, $lte: 2015 } } )
```

Selection

Query parameter describes the documents we are interested in



Conditions on fields (each field can be used **at most once**)

1 Value equality

- The actual field value must be **identical to** the **specified value** (including, e.g., the order of nested fields or array items)

2 Query operators (each operator can be used **at most once**)

- The actual field value must **satisfy all** the provided operators

Value Equality Conditions: Examples

Select movies having a specific director

```
1 db.movies.find(  
2   { director: { firstname: "David", lastname: "Russell" } }  
3 )
```

```
1 db.movies.find(  
2   { director: { lastname: "Russell", firstname: "David" } }  
3 )
```

Select movies having specific actors


```
1 db.movies.find( { actors: [ ObjectId("7"), ObjectId("5") ] } )
```

```
1 db.movies.find( { actors: [ ObjectId("5"), ObjectId("7") ] } )
```

Queries in both the pairs are not equivalent!

Dot Notation

The dot notation for field names



1 Accessing fields of embedded documents

- "field.subfield"
 - e.g.: "director.firstname"

2 Accessing items of arrays

- "field.index"
 - e.g.: "actors.2"
 - Positions start at 0

Value Equality Conditions

Example (revisited)

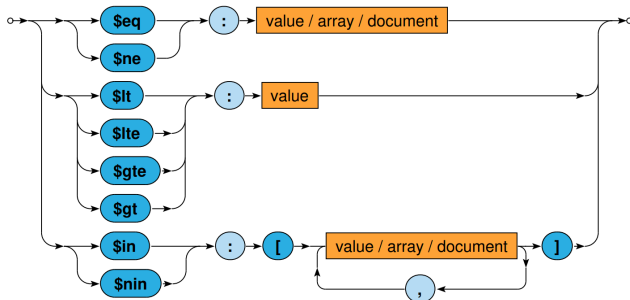
Select movies having a specific director

```
1 db.movies.find(  
2   { director: { firstname: "David", lastname: "Russell" } }  
3 )
```

```
1 db.movies.find(  
2   { "director.firstname": "David", "director.lastname": "Russell" }  
3 )
```

Query Operators

Comparison operators



Comparisons take particular **BSON** data types into account

- Certain numeric conversions are automatically applied

Query Operators

Comparison operators

1 \$eq, \$ne

- Tests the actual field value for **equality/inequality**
 - The same behavior as in case of value equality conditions

2 \$lt, \$lte, \$gte, \$gt

- Tests whether the actual field value **is** less than/less than or equal/greater than or equal/greater than the provided value

3 \$in

- Tests whether the actual field value is **equal to** at least one of the provided values


4 \$nin

- Negation of \$in

Query Operators

Element operators

- 1 **\$exists** – tests whether a given field **exists/not exists**




Evaluation operators

- 1 **\$regex** – tests whether a given field value **matches** a specified regular expression (PCRE)
- 2 **\$text** – **performs** text search (text index must exist)

Query Operators

Array operators

- 1 **\$all** – tests whether a given array **contains all** the specified items (in any order)



Example (revisited)

Select movies having specific actors

```

1 db.movies.find(
2   { actors: [ ObjectId("5"), ObjectId("7") ] }
3 )
  
```

```

1 db.movies.find(
2   { actors: { $all: [ ObjectId("5"), ObjectId("7") ] } }
3 )
  
```

Query Operators

Array operators

- 1 **\$size** – **tests the size** of a given array against a fixed number (and not, e.g., a range, unfortunately)




- 2 **\$elemMatch** – tests whether a given array **contains** at least one item that satisfies all the involved query operations



Query Operators

Logical operators

1 \$and, \$or



- Logical connectives for **conjunction/disjunction**
- **At least two** involved query expressions must be provided
- Only allowed at the **top level** of a query

2 \$not



- Logical **negation of** exactly one involved query operator
- i.e. **cannot** be used at the top level of a query

Querying Arrays

Condition based on **value equality** is satisfied when...

- 1 the given **field** as a **whole** is **identical** to the provided value, or
- 2 at **least one item** of the **array** is **identical** to the provided value

```
1 db.movies.find( { actors: ObjectId("5") } )
```

```
1 { actors: ObjectId("5") }
```

```
1 { actors: [ ObjectId("5"), ObjectId("7") ] }
```


Querying Arrays

Condition based on **query operators** is satisfied when...

- 1 the given **field** as a **whole** **satisfies** all the involved operators, or
- 2 **at least one** item of the array **satisfies** each of the involved operators
 - note, however, that this item might **not** be the **same** for all the **individual operators**

```
1 db.movies.find( { ratings: { $gte: 2, $lte: 3 } } )
```


```
1 { ratings: 3 }
```

```
1 { ratings: [ 3, 7, 5 ] }
```

Use `$elemMatch` when just one array item should be found for all the operators

Projection

Projection allows us to determine the **fields returned** in the result



- **true** or **1** for fields to be **included**
- **false** or **0** for fields to be **excluded**
- **Positive** and **negative** enumerations **cannot** be **combined**!
 - The only exception is `_id` which is included by default
- **Projection operators** – allow to **select** particular array items


Projection Operators

Array operators

- 1 **\$elemMatch** – selects the **first matching item** of an array
 - This item must **satisfy all** the operators included in query
 - When there is **no such item**, the field is **not returned** at all



- 2 **\$slice** – selects the **first count items** of an array (when count is **positive**)/the **last count items** (when **negative**)
 - Certain number of items can also be skipped



Projection: Examples

Find a particular movie, select its identifier, title and actors

```

1 db.movies.find(
2   { _id: ObjectId("2") },
3   { title: true, actors: true }
4 )

```

```

1 {
2   _id: ObjectId("2"),
3   title: "Zombieland",
4   actors: [ ObjectId("6"),
5             ObjectId("4"),
6             ObjectId("5") ]
7 }

```

Find movies from 2000, select their titles and the last two actors

```

1 db.movies.find(
2   { year: 2000 },
3   {
4     title: 1, _id: 0,
5     actors: { $slice: -2 }
6   }
7 )

```

```

1 {
2   title: "Zombieland",
3   actors: [ ObjectId("4"),
4             ObjectId("5") ]
5 }

```

Modifiers

Modifiers change the **order** and **number** of returned documents

- 1 **sort** – orders the documents in the result
- 2 **limit** – returns at most a certain number of documents




- 3 **skip** – skips a certain number of documents from the beginning



All the modifiers are **optional**, can be chained in **any order**, but must all be specified **before** any documents are retrieved via a given cursor

Modifiers

Sort modifier orders the documents in the result



- **1** for ascending, **-1** for descending order
- The order of documents is **undefined** unless **explicitly sorted**
- Sorting of **larger datasets** should be supported by **indices**
- Sorting happens **before** the **projection phase**
 - i.e. **not included fields** can be used for sorting purposes as well

MongoDB Features

- Document-Oriented storage
- Full Index Support
- Replication High Availability
- Auto-Sharding
- Querying
- Fast In-Place Updates
- Map/Reduce functionality

Lecture Conclusion

MongoDB

- Document database for JSON documents
- Sharding with master-slave replication architecture

Query functionality

- CRUD operations
 - Insert, find, update, remove
 - Complex filtering conditions
- MapReduce
- Index structures

Index Functionality

- **B+ tree** indexes
- An index is **automatically created** on the `_id` field (the primary key)
- Users can create other indexes to improve query performance or to enforce Unique values for a particular field
- Supports single field index as well as **Compound index**
 - Like SQL order of the fields in a compound index matters
 - If you index a field that holds an array value, MongoDB creates separate index entries for every element of the array
- **Sparse** property of an index ensures that the index only contain entries for documents that have the indexed field. (so ignore records that do not have the field defined)
- If an index is both **unique** and **sparse** – then the system will reject records that have a duplicate key value but allow records that do not have the indexed field defined

Aggregated functionality

Aggregation framework provides SQL-like aggregation functionality

- Pipeline documents from a collection pass through an aggregation pipeline, which transforms these objects as they pass through
- Expressions produce output documents based on calculations performed on input documents

```
1 db.parts.aggregate(  
2   {$group : {  
3     _id: type, totalquantity: { $sum: quantity}  
4   }  
5 }  
6 )
```

Map reduce functionality

- Performs complex aggregator functions given a collection of keys, value pairs
- Must provide at least a map function, reduction function and a name of the result set

```
1 db.collection.mapReduce(  
2   <mapfunction>, <reducefunction>, {  
3     out: <collection>,  
4     query: <document>,  
5     sort: <document>,  
6     limit: <number>,  
7     finalize: <function>,  
8     scope: <document>,  
9     jsMode: <boolean>,  
10    verbose: <boolean>  
11  }  
12 )
```

Indexes: High performance read

- Typically used for frequently used queries
- Necessary when the total size of the documents exceeds the amount of available RAM.
- Defined on the collection level
 - Can be defined on 1 or more fields
 - Composite index (SQL) \Rightarrow Compound index (MongoDB)
- B-tree index
- Only 1 index can be used by the query optimizer when retrieving data
- Index covers a query - match the query conditions and return the results using only the index;
 - Use index to provide the results.

Replication of data

Ensures redundancy, backup, and automatic failover

- Recovery manager in the RDMS

Replication occurs through groups of servers known as replica sets

- Primary set – set of servers that client tasks direct updates to
- Secondary set – set of servers used for duplication of data
- At the most can have 12 replica sets
 - Many different properties can be associated with a secondary set i.e. secondary-only, hidden delayed, arbiters, non-voting
- If the primary set fails the secondary sets 'vote' to elect the new primary set

Consistency of data

All read operations issued to the primary of a replica set are consistent with the last write operation

- Reads to a primary have strict consistency
 - Reads reflect the latest changes to the data
- Reads to a secondary have eventual consistency
 - Updates propagate gradually
- If clients permit reads from secondary sets – then client may read a previous state of the database
- Failure occurs before the secondary nodes are updated
 - System identifies when a rollback needs to occur
 - Users are responsible for manually applying rollback changes

Provides Memory Mapped Files

- A memory-mapped file is a segment of virtual memory which has been assigned a direct byte-for-byte correlation with some portion of a file or file-like resource.
- `mmap()`

Other additional features

Supports geospatial data of type

- Spherical
 - Provides longitude and latitude
- Flat
 - 2 dimensional points on a plane
- Geospatial indexes